

Data-Driven Decision Trees for Business

Felipe O. Cerezer

2024-08-25

Step 1: Required Libraries

We will use the following R packages: - `rpart` for building decision trees - `caret` for various modeling techniques, including decision trees - `ggplot2` for data visualizations

```
## Install required packages if not already installed
#install.packages(c("rpart", "caret", "ggplot2"))
```

```
## Load libraries
library(rpart)
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(ggplot2)
```

Step 2: Generate Fake Data

First, I will create a synthetic dataset representing customers information, such as Age, Income, spending habits score, and Segment (the target variable).

```
# Set seed for reproducibility
set.seed(42)
```

```
# Generate fake data
n <- 200 # Number of observations
data <- data.frame(
  Age = sample(18:70, n, replace = TRUE),
  Income = sample(20000:120000, n, replace = TRUE),
  SpendingScore = sample(1:100, n, replace = TRUE),
  Segment = sample(c("Low", "Medium", "High"), n, replace = TRUE)
)
```

```
# View the first few rows of the data
head(data)
```

```
##   Age Income SpendingScore Segment
## 1  66 117123           75      Low
## 2  54  95338           13      High
## 3  18  81640           18   Medium
## 4  42  32349           85      Low
## 5  27  54782            5      High
```

```
## 6 53 55376          31    High
```

Explanation: Here, I used the `sample()` function to randomly generate values for Age, Income, and SpendingScore. The Segment is a categorical variable indicating whether a customer belongs to the “Low”, “Medium”, or “High” segment.

Step 3: Splitting the Data

To evaluate the decision tree’s performance, I split the dataset into training (70%) and testing (30%) sets. The model will be trained on the training set and tested on the testing set to check its accuracy.

```
# Create a training (70%) and testing (30%) split
set.seed(42)
trainIndex <- createDataPartition(data$Segment, p = 0.7, list = FALSE)
trainData <- data[trainIndex, ]
testData <- data[-trainIndex, ]
```

Explanation: The `createDataPartition()` function from the `caret` package creates an index to split the data into training and testing sets. I used a 70/30 split, which is a common practice to ensure the model has enough data to learn from while still being tested on unseen data.

Step 4: Building the Decision Tree

Next, I use the `rpart()` function to build the decision tree model. This function will automatically determine the best splits in the data based on the Segment variable.

```
# Build the decision tree model
treeModel <- rpart(Segment ~ Age + Income + SpendingScore,
                   data = trainData,
                   method = "class")

# Print the model summary
summary(treeModel)
```

```
## Call:
## rpart(formula = Segment ~ Age + Income + SpendingScore, data = trainData,
##       method = "class")
##      n= 141
##
##           CP nsplit rel error  xerror      xstd
## 1 0.11956522     0 1.0000000 1.228261 0.05148976
## 2 0.04891304     1 0.8804348 1.076087 0.05902620
## 3 0.02173913     3 0.7826087 1.097826 0.05818262
## 4 0.01630435     7 0.6956522 1.086957 0.05861308
## 5 0.01000000     9 0.6630435 1.097826 0.05818262
##
## Variable importance
## SpendingScore      Income      Age
##           60           26           14
##
## Node number 1: 141 observations,      complexity param=0.1195652
##   predicted class=Medium expected loss=0.6524823 P(node) =1
##   class counts:      47      45      49
##   probabilities: 0.333 0.319 0.348
##   left son=2 (39 obs) right son=3 (102 obs)
```

```

## Primary splits:
##   Income      < 45830.5 to the left,  improve=2.6506530, (0 missing)
##   SpendingScore < 67.5   to the left,  improve=1.8436950, (0 missing)
##   Age          < 20.5   to the right, improve=0.7159897, (0 missing)
## Surrogate splits:
##   SpendingScore < 98.5   to the right, agree=0.738, adj=0.051, (0 split)
##
## Node number 2: 39 observations,    complexity param=0.01630435
## predicted class=Low    expected loss=0.5128205 P(node) =0.2765957
##   class counts:      12    19    8
##   probabilities: 0.308 0.487 0.205
## left son=4 (31 obs) right son=5 (8 obs)
## Primary splits:
##   SpendingScore < 80      to the left,  improve=2.5312240, (0 missing)
##   Income        < 24848   to the right, improve=0.7763278, (0 missing)
##   Age           < 40.5   to the right, improve=0.7065527, (0 missing)
##
## Node number 3: 102 observations,    complexity param=0.04891304
## predicted class=Medium expected loss=0.5980392 P(node) =0.7234043
##   class counts:      35    26    41
##   probabilities: 0.343 0.255 0.402
## left son=6 (72 obs) right son=7 (30 obs)
## Primary splits:
##   SpendingScore < 67.5   to the left,  improve=1.732353, (0 missing)
##   Income        < 58294   to the left,  improve=1.671242, (0 missing)
##   Age           < 37.5   to the left,  improve=1.368311, (0 missing)
## Surrogate splits:
##   Age < 68      to the left,  agree=0.725, adj=0.067, (0 split)
##
## Node number 4: 31 observations,    complexity param=0.01630435
## predicted class=High    expected loss=0.6129032 P(node) =0.2198582
##   class counts:      12    12    7
##   probabilities: 0.387 0.387 0.226
## left son=8 (17 obs) right son=9 (14 obs)
## Primary splits:
##   SpendingScore < 41.5   to the left,  improve=1.2466790, (0 missing)
##   Age           < 39.5   to the right, improve=0.9876181, (0 missing)
##   Income        < 35452   to the left,  improve=0.7700579, (0 missing)
## Surrogate splits:
##   Income < 25799.5 to the right, agree=0.710, adj=0.357, (0 split)
##   Age < 34.5      to the right, agree=0.645, adj=0.214, (0 split)
##
## Node number 5: 8 observations
## predicted class=Low    expected loss=0.125 P(node) =0.05673759
##   class counts:      0    7    1
##   probabilities: 0.000 0.875 0.125
##
## Node number 6: 72 observations,    complexity param=0.04891304
## predicted class=High    expected loss=0.625 P(node) =0.5106383
##   class counts:      27    21    24
##   probabilities: 0.375 0.292 0.333
## left son=12 (59 obs) right son=13 (13 obs)
## Primary splits:
##   Age           < 56.5   to the left,  improve=2.214146, (0 missing)

```

```

##      SpendingScore < 61.5    to the right, improve=1.569780, (0 missing)
##      Income          < 58643  to the left,  improve=1.337302, (0 missing)
##
## Node number 7: 30 observations
##   predicted class=Medium   expected loss=0.4333333   P(node) =0.212766
##   class counts:      8      5      17
##   probabilities: 0.267 0.167 0.567
##
## Node number 8: 17 observations
##   predicted class=High     expected loss=0.4705882   P(node) =0.1205674
##   class counts:      9      6      2
##   probabilities: 0.529 0.353 0.118
##
## Node number 9: 14 observations
##   predicted class=Low      expected loss=0.5714286   P(node) =0.09929078
##   class counts:      3      6      5
##   probabilities: 0.214 0.429 0.357
##
## Node number 12: 59 observations,    complexity param=0.02173913
##   predicted class=High     expected loss=0.559322   P(node) =0.4184397
##   class counts:      26     16     17
##   probabilities: 0.441 0.271 0.288
##   left son=24 (8 obs) right son=25 (51 obs)
##   Primary splits:
##       SpendingScore < 58.5    to the right, improve=1.579595, (0 missing)
##       Income          < 108279 to the left,  improve=1.414933, (0 missing)
##       Age              < 20.5   to the right, improve=1.057832, (0 missing)
##
## Node number 13: 13 observations
##   predicted class=Medium   expected loss=0.4615385   P(node) =0.09219858
##   class counts:      1      5      7
##   probabilities: 0.077 0.385 0.538
##
## Node number 24: 8 observations
##   predicted class=High     expected loss=0.25   P(node) =0.05673759
##   class counts:      6      0      2
##   probabilities: 0.750 0.000 0.250
##
## Node number 25: 51 observations,    complexity param=0.02173913
##   predicted class=High     expected loss=0.6078431   P(node) =0.3617021
##   class counts:      20     16     15
##   probabilities: 0.392 0.314 0.294
##   left son=50 (9 obs) right son=51 (42 obs)
##   Primary splits:
##       SpendingScore < 49.5    to the right, improve=2.0429510, (0 missing)
##       Income          < 58643  to the right, improve=1.4464200, (0 missing)
##       Age              < 20.5   to the right, improve=0.7579577, (0 missing)
##
## Node number 50: 9 observations
##   predicted class=Low      expected loss=0.3333333   P(node) =0.06382979
##   class counts:      2      6      1
##   probabilities: 0.222 0.667 0.111
##
## Node number 51: 42 observations,    complexity param=0.02173913

```

```

## predicted class=High expected loss=0.5714286 P(node) =0.2978723
## class counts: 18 10 14
## probabilities: 0.429 0.238 0.333
## left son=102 (25 obs) right son=103 (17 obs)
## Primary splits:
## SpendingScore < 25.5 to the left, improve=2.4475070, (0 missing)
## Income < 56573 to the right, improve=0.9523810, (0 missing)
## Age < 44.5 to the right, improve=0.3150183, (0 missing)
## Surrogate splits:
## Income < 109366 to the left, agree=0.690, adj=0.235, (0 split)
## Age < 35.5 to the left, agree=0.643, adj=0.118, (0 split)
##
## Node number 102: 25 observations, complexity param=0.02173913
## predicted class=Low expected loss=0.6 P(node) =0.177305
## class counts: 9 10 6
## probabilities: 0.360 0.400 0.240
## left son=204 (10 obs) right son=205 (15 obs)
## Primary splits:
## Income < 82119.5 to the right, improve=1.4533330, (0 missing)
## Age < 34 to the right, improve=0.8712821, (0 missing)
## SpendingScore < 15.5 to the right, improve=0.5200000, (0 missing)
## Surrogate splits:
## SpendingScore < 3.5 to the left, agree=0.68, adj=0.2, (0 split)
##
## Node number 103: 17 observations
## predicted class=High expected loss=0.4705882 P(node) =0.1205674
## class counts: 9 0 8
## probabilities: 0.529 0.000 0.471
##
## Node number 204: 10 observations
## predicted class=High expected loss=0.4 P(node) =0.07092199
## class counts: 6 3 1
## probabilities: 0.600 0.300 0.100
##
## Node number 205: 15 observations
## predicted class=Low expected loss=0.5333333 P(node) =0.106383
## class counts: 3 7 5
## probabilities: 0.200 0.467 0.333

```

Explanation: The formula `Segment ~ Age + Income + SpendingScore` indicates that I predicted the Segment variable based on Age, Income, and SpendingScore. The method = “class” argument specifies that this is a classification problem.

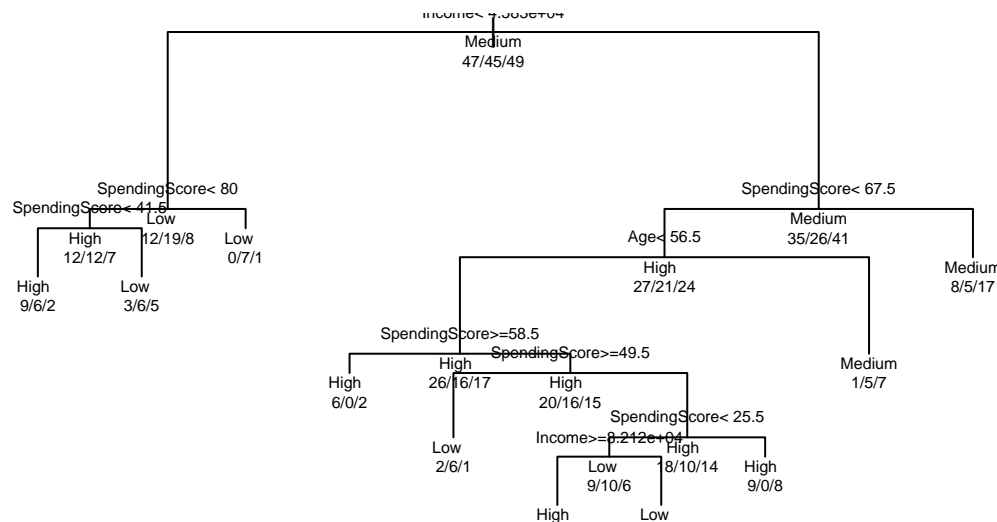
Step 5: Visualizing the Decision Tree

Now, I will visualize the decision tree, which helps us understand the decision rules the model has learned.

```

# Plot the decision tree
plot(treeModel)
text(treeModel, use.n = TRUE, all = TRUE, cex = 0.5)

```



Explanation: The `plot()` function generates a visual representation of the decision tree, while `text()` adds labels to the tree nodes. The `use.n = TRUE` argument includes the number of observations in each node, and `cex = 0.8` controls the size of the text. The plot typically shows the decision tree with nodes and branches. Each node represents a decision rule (e.g., “Is Age < 40?”), and branches represent the outcomes of these decisions.

Step 6: Predicting and Evaluating the Model

After training the model, I use it to make predictions on the testing set and evaluate its accuracy using a confusion matrix.

```
# Predict on the test data
predictions <- predict(treeModel, newdata = testData, type = "class")

#Ensure the reference variable is a factor
testData$Segment <- factor(testData$Segment)

# Predict on test data
predictions <- factor(predictions, levels = levels(testData$Segment))

# Generate the confusion matrix
confusionMatrix(predictions, testData$Segment)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction High Low Medium
##      High      9   5     7
##      Low       2   5     5
##      Medium    8   9     9
##
## Overall Statistics
##
##               Accuracy : 0.3898
##               95% CI : (0.2655, 0.5256)
##      No Information Rate : 0.3559
##      P-Value [Acc > NIR] : 0.3379
```

```
##
##           Kappa : 0.0797
##
## Mcnemar's Test P-Value : 0.4762
##
## Statistics by Class:
##
##           Class: High Class: Low Class: Medium
## Sensitivity           0.4737    0.26316    0.4286
## Specificity           0.7000    0.82500    0.5526
## Pos Pred Value        0.4286    0.41667    0.3462
## Neg Pred Value        0.7368    0.70213    0.6364
## Prevalence            0.3220    0.32203    0.3559
## Detection Rate        0.1525    0.08475    0.1525
## Detection Prevalence  0.3559    0.20339    0.4407
## Balanced Accuracy     0.5868    0.54408    0.4906
```

Explanation: The `predict()` function applies the trained model to the test data to generate predictions. The `confusionMatrix()` function from the `caret` package compares the predictions to the actual segments, allowing us to evaluate the model's accuracy.

Step 7: Fine-Tuning the Model

To improve the model's performance, I can use the `caret` package to fine-tune the decision tree by cross-validating different complexity parameters.

```
# Define the training control
train_control <- trainControl(method = "cv", number = 10)

# Train the model with caret, applying cross-validation
tunedModel <- train(Segment ~ Age + Income + SpendingScore,
                    data = trainData,
                    method = "rpart",
                    trControl = train_control)

# Print the best model
print(tunedModel$finalModel)
```

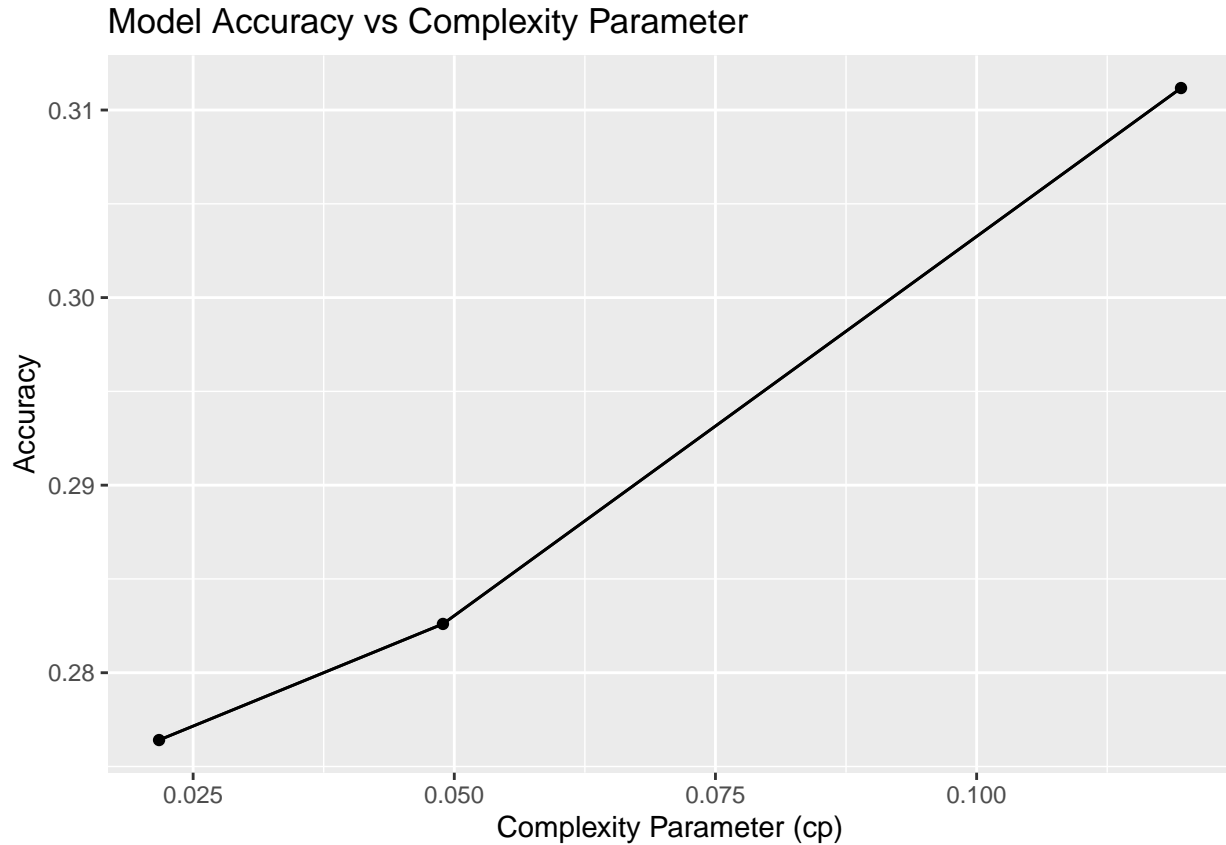
```
## n= 141
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 141 92 Medium (0.3333333 0.3191489 0.3475177) *
```

Explanation: The `trainControl()` function specifies that we want to use 10-fold cross-validation (`method = "cv"`, `number = 10`). The `train()` function from the `caret` package then trains the model, tuning it by testing different values of the complexity parameter (`cp`).

Step 8: Visualization

Finally, I can visualize the performance of our tuned model using `ggplot2`, which allows us to plot the relationship between the complexity parameter and accuracy.

```
# Plot the complexity parameter (cp) vs accuracy  
ggplot(tunedModel) +  
  geom_line(aes(x = tunedModel$results$cp, y = tunedModel$results$Accuracy)) +  
  labs(x = "Complexity Parameter (cp)", y = "Accuracy") +  
  ggtitle("Model Accuracy vs Complexity Parameter")
```



Explanation: Here, `ggplot()` initializes the plot, and `geom_line()` creates a line graph showing how the accuracy of the model varies with the complexity parameter (cp).